

Does draft-order modeling beat expert consensus?

A counterfactual evaluation of five drafting policies

Yousuf Rashid
University of Waterloo

Abstract

I built a system to pick fantasy basketball players better than the default rankings do. It has three parts: an XGBoost projection model, a Cox proportional-hazards model for when each player comes off the board, and a Monte Carlo layer that uses the second to decide how to spend the first. Then I tested the whole thing against the ranking it was supposed to beat, over 4,400 replayed drafts.

It lost. Average draft position (ADP) predicts realized production better than anything I fit ($\rho = 0.574$ against 0.555 for my value estimate and 0.529 after a positional adjustment), and every policy built on those estimates finished behind the policy that just follows ADP. The survival model itself is accurate, with concordance 0.935 and expected calibration error 0.029. Knowing when a player comes off the board just does not improve the pick.

The interesting part is how the answer changed with sample size. At 132 paired drafts the sequencing policy looked like it won by +2.35 points. At 1,056 the edge was -0.14 and it lost 54% of head-to-head drafts. The first number was noise, and I would have reported it if I had stopped there.

1 What the system does

A snake draft has a structure that rewards thinking ahead. You pick at 6, then not again until 19, then 30. Thirteen players leave the board between your turns. If you want two of them, the order you take them in matters: take the one who will still be there later, and you lose him.

That is the whole idea. Formally, at pick j with a next turn at k , the value of waiting on player i depends on

$$P(\text{available at } k \mid \text{available at } j) = \frac{S_i(k)}{S_i(j)},$$

the conditional survival ratio rather than $S_i(k)$ on its own. The distinction matters more than it looks. A player with an ADP of 5 who is somehow still on the board at pick 20 is a different proposition from the same player before the draft started, and only the ratio knows that.

Three models feed the decision.

Projections (Model A). XGBoost per category, trained on season-to-season transitions with lag features, age, and team-context terms. Rookies have no history to lag, so they shrink toward a prior fit on what players from the same draft-pick tier actually did, weighted $w = n/(n + k)$ on career minutes.

Availability (Model B). A Cox proportional-hazards model over draft durations, with Kaplan–Meier curves as the non-parametric fallback. Held-out concordance is 0.935; expected calibration error is 0.029, which matters because the decision layer consumes the probability, not the ranking.

Games played. Added after the first backtest, and the one place a model clearly helped. The projection code multiplied every per-minute rate by a flat 82 games, so it assumed everyone plays a full season. League-average games played is about 46. A games-played model cut RMSE from 38.3 to 18.5, a 52% reduction against the flat assumption and 19% against repeating last season’s total.

2 The backtest

This is counterfactual policy evaluation: estimating what a policy would have earned on data generated by a different policy. My agent drafts from one slot using the policy under test. The other teams draft from that season’s actual ADP board, jittered per seed so the opponents are not identical every run. Rosters are scored on *realized* end-of-season production, not on projections. Score on projections and a policy can win by agreeing with its own model.

Table 1: Sweep dimensions. Every comparison is paired on (season, league size, slot, seed), so the draft’s own luck cancels.

Seasons	2020–21 through 2025–26 (6)
League sizes	10 and 12 teams
Draft slots	all, exhaustively
Seeds	8 per configuration
Rollout simulations	60 per candidate
Total drafts	4,400

Roster legality lives inside the rollout rather than in a separate optimizer. Starting slots (4G/4F/2C) fill before the three bench seats, so the bench is the slack that absorbs a positional run. A draft that returns 12 or 14 players is a bug, and the harness asserts against it.

Two details cost me real time. Simulating to the end of the draft was wasted work, since picks after your last one cannot change your roster; stopping there was the single largest speedup. And scoring every candidate against *common random numbers*, meaning the same opponent draws, made the argmax stable at 60 simulations instead of the 1,000 I had budgeted. I checked 30 and it was not enough: only one of three draft states gave a stable answer across seeds.

3 The ablation ladder

Five policies, one harness. Rungs 0 through 3 see identical inputs, so any difference between them is the decision rule and nothing else. Only rung 4 swaps in my own projections, which is what makes rung 4 minus rung 3 a separate measurement of whether the projection model was worth anything.

Rung 3 is the comparison the design exists for. It holds projections fixed against rung 2 and changes only whether the policy simulates forward, so it isolates sequencing from projection

Table 2: Realized starter value, 1,056 paired drafts per rung (176 for rung 4, which needs projections and so runs on 2025–26 only). Win rate is against the rung below.

Rung	Policy	Mean	Edge	Win rate	90% interval
0	Follow ADP	61.50	–	–	–
1	Best available by value	60.21	-1.29	0.375	$[-14.0, +10.7]$
2	VORP, greedy	54.22	-5.99	0.293	$[-26.2, +14.7]$
3	VORP + sequencing	54.07	-0.14	0.457	$[-14.2, +14.6]$
4	Own projections + both	3.76	-46.32	0.023	$[-73.7, -11.0]$

quality. The answer is -0.14 points and a 45.7% win rate. Sequencing does not help. Figure 1 shows why the means alone understate how little separates the rungs: the distributions overlap almost entirely.

Against ADP directly, rung 3 wins 25.6% of paired drafts and zero of twelve draft slots. It is not that sequencing works from some positions and not others. It loses everywhere. Figure 2 is the same comparison per draft rather than in aggregate.

3.1 The result I nearly published

My first run used one season, two league sizes, and six seeds: 132 paired drafts. Rung 3 beat rung 2 by $+2.35$ with a 55.3% win rate, and the per-combination numbers pointed the same way both times. I wrote it down as “directionally consistent.”

It was not. The 90% interval on that estimate was $[-16.3, +25.3]$, roughly ten times wider than the effect. At eight times the sample the edge collapsed to -0.14 . Nothing about the model changed between those two runs.

I am putting this in the paper because the interval was visible the whole time and I still described a coin flip as a direction. If the sweep had stopped at 132 drafts, the writeup would have claimed a win.

4 Why every learned policy lost

A negative result is only useful with a diagnosis, so I tested each value estimate directly against realized production. Rank correlation, six seasons, higher is better.

Table 3: How well each ranker predicts realized value (Spearman ρ).

Ranker	Full board	Top 60
ADP rank	0.574	0.368
Board-implied value	0.555	0.360
Value over replacement	0.529	0.316
Own projections	0.331	0.270

ADP wins on the full board and again in the top 60, where drafts are actually decided, and Figures 3 and 4 show the consequence draft by draft and slot by slot. So rungs 1 through 4 are not losing because the decision layer is broken. They are optimizing a worse objective than the

ranking they are trying to beat, and no amount of sequencing fixes an inferior input. Rung 3 was doing competent inference on the wrong numbers.

This is a market-efficiency finding, and in hindsight it should have been the prior. ADP is a consensus of experts who update daily on injury news, minutes expectations, and beat reporting. My projections see prior-season box scores. When those two disagree, the box scores are usually the ones missing something.

4.1 Where the result does and does not hold

The finding survives both cuts of the sweep. Splitting by season checks that no single year is carrying it, and splitting by league size checks the mechanism: a deeper board means longer waits between picks, so if sequencing worked at all it should work best in the smallest league.

4.2 One correction to the table

Rung 2's -5.99 overstates its failure, and the cause is my scoring. The ladder above scores starters only. VORP-greedy drafts depth on purpose and strands 9.23 points on the bench, against 5.05 for ADP. Scored on total roster value the penalty falls to -2.10 and its win rate rises from 29% to 44%.

Rung 3 stays negative under both metrics (-0.14 starters, -2.48 total), so the sequencing conclusion holds either way. But the honest statement about rung 2 is that a chunk of its deficit was a choice I made about scoring, not something the policy did.

4.3 What the projection model was worth

Rung 4 is the weakest rung at 3.76, and most of that traces to one thing the data cannot supply. Four players (VanVleet, Lillard, Irving, Haliburton) were projected for 61 to 66 games and played zero. All four were ruled out after the ADP board published.

No feature built from prior-season statistics can know that. A live injury feed could, and the database has no historical one: `player_status` is a single snapshot with no per-season history, so a past draft's injury designations cannot be reconstructed. Reading today's status onto a 2018 row would leak the future, since today's status partly reflects what happened after that draft.

So the games-played model corrects the population-level bias, which was real and large, and cannot correct individual unforeseeable injuries. I am not going to pretend otherwise by fitting something that appears to predict them. The ADP board handles these cases and my model does not, and that gap is a large part of why rung 4 finishes where it does.

5 A data bug worth recording

Partway through the diagnosis, VORP-greedy was drafting players with high value and zero realized production. Four of them on the 2025–26 board had no position at all.

The ADP ingest normalized names by stripping generational suffixes before matching, so “Jaren Jackson Jr.” collapsed onto “Jaren Jackson.” The lookup kept the lowest player id, which is the father. Seven players were linked to their fathers across all twelve seasons: Jackson, Hardaway, Nance, Porter, Trent, Smith, Payton. The fathers have no current roster row and no

current statistics, so those players became position-less entries worth nothing, which is worse than being unmatched because they still carried an inflated VORP.

The fix matches on the exact name first and falls back to the suffix-stripped form only when no player claims the exact one. Suffix-stripping still has to exist, because “Russell Westbrook III” and “Russell Westbrook” are the same man, so the fallback map refuses to answer when two real players share a base name. Seven links corrected, no regressions, and the same five unmatched names as before, all of whom never played an NBA game.

This is the same failure mode the project had already hit and fixed for accented names, where Jokić and Dončić silently failed to link. Suffixes were the case that got missed.

6 What I would do next

Not more sequencing. The survival model works and the decision layer works; the inputs are the problem, and the diagnosis says ADP already contains what my projections estimate. Beating it needs signal ADP does not have. A live injury feed is the obvious candidate, since that is the one gap I can point at with a number attached.

The part I would keep is the harness. Four thousand paired off-policy drafts with per-slot and per-seed distributions is the machinery that told me the +2.35 was noise, and it would have told me the same thing about any other idea I tried.

Reproducing this

- `models/projection/` — XGBoost projections, games-played model, rookie priors, walk-forward validation.
- `models/behaviour/` — Kaplan–Meier and Cox survival models, calibration, and the availability query the rollout calls.
- `models/decision/` — rollout, the five policies, the backtest harness, plots.
- `artifacts/backtest.csv` — all 4,400 drafts, one row each.
- `artifacts/ablation_ladder.csv` — the ladder table.
- `tests/` — 141 tests. The 70 that need no database run in CI on every push; the rest skip, since the database is published as a release asset rather than committed.

Every module also runs standalone as a worked example. Splits are temporal everywhere: a draft is a point in time, and training on 2025 to predict 2019 would be a leak that reports a better number for itself. Leak detection is a gate rather than only a test — the rate trainer refuses to run if it fires, and beyond checking feature names it permutes the current season’s statistics and rebuilds, since a feature whose value moves was reading the season it is meant to predict.




Figure 1: The ladder as distributions rather than means. The violins are realized starter value across all 1,056 paired drafts per rung; the lower panel is the paired difference against the rung below with its 5th–95th percentile band. Every band except rung 4’s contains zero, which is the shape of the whole result: the losses are real but small against the draft’s own variance.




Figure 2: Sequencing on identical inputs, draft by draft. If simulating the board forward helped, this would sit to the right of zero. It is centred on it, with a median of -0.44 and a 5th–95th range of $[-14.2, +14.6]$: the per-draft spread is two orders of magnitude wider than the effect being measured.




Figure 3: Share of paired drafts each policy wins against the incumbent. A policy that merely had higher variance would sit near 50%; every rung sits below it.




Figure 4: Edge over ADP by draft slot. Sequencing should pay most from the turn ends, where the wait between picks is longest, and it does not pay anywhere: rung 3 loses at all twelve slots. A policy that worked from some positions and not others would show a slope here.




Figure 5: Mean realized starter value by season. The rung ordering is stable across all six, so the result is not one season's injuries. Absolute values move a lot between seasons, which is why every comparison in this paper is paired.




Figure 6: Edge over ADP by league size. Ten and twelve teams give the same answer. The mechanism predicted a gradient here and there is not one.